

XGBoost

December 30, 2025

1 XGBoost

2 Project: Fraud Detection - XGBoost Model

Mục tiêu: Tối ưu hóa sai số (Bias) và đạt độ chính xác cao nhất (SOTA Performance).

2.1 Tại sao XGBoost?

- **Gradient Boosting:** Học sửa sai từ các cây trước đó, giúp tối ưu hóa các trường hợp khó phân loại (Hard examples).
- **Regularization:** Tích hợp L1/L2 regularization giúp mô hình tổng quát hóa tốt hơn Random Forest.
- **Handling Imbalance:** Tham số `scale_pos_weight` hỗ trợ tốt cho dữ liệu lệch.

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')

# Thư viện scikit-learn
from sklearn.model_selection import train_test_split, StratifiedKFold, cross_val_score
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.metrics import (
    classification_report, confusion_matrix, roc_auc_score,
    average_precision_score, precision_recall_curve,
    precision_score, recall_score, f1_score, roc_curve
)

# Thư viện xử lý mất cân bằng
from imblearn.over_sampling import SMOTE, RandomOverSampler
from imblearn.under_sampling import RandomUnderSampler
from imblearn.combine import SMOTEENN, SMOTETomek
from imblearn.pipeline import Pipeline as ImbPipeline
```

```
# Thư viện XGBoost và tối ưu
import xgboost as xgb
import optuna
import joblib
import os
```

3 LOAD DỮ LIỆU VÀ TIỀN XỬ LÝ

```
[2]: # Load dữ liệu
df = pd.read_csv("C://Users/hi//Downloads//fraud_detection_project//data//
↳processed//fraud_feature_engineering.csv")
df
```

```
[2]:
```

	Unnamed: 0	step	type	amount	nameOrig	oldbalanceOrig	\
0	0	401	CASH_IN	290670.88	C1172770329	30613326.59	
1	1	346	CASH_OUT	74969.02	C361253330	0.00	
2	2	34	CASH_OUT	152359.15	C1296741848	0.00	
3	3	302	TRANSFER	553238.82	C765587136	0.00	
4	4	130	CASH_IN	432624.19	C852698921	18910.00	
...	...	...	...	...	...	...	
325929	325929	282	CASH_OUT	34345.29	C943859349	20380.00	
325930	325930	379	PAYMENT	8234.77	C1729134111	272.00	
325931	325931	235	CASH_OUT	13499.45	C297276080	0.00	
325932	325932	369	CASH_OUT	236678.98	C2094429825	10018.00	
325933	325933	395	CASH_OUT	108989.28	C1625687711	337544.00	
...	...	...	...	...	...	...	
	newbalanceOrig	nameDest	oldbalanceDest	newbalanceDest	...	\	
0	30903997.48	C1360884032	6252769.10	5962098.22	...		
1	0.00	C1628643300	1510659.40	1585628.42	...		
2	0.00	C1234551068	519498.85	671857.99	...		
3	0.00	C845383981	1278572.18	1831811.01	...		
4	451534.19	C290472837	306529.18	0.00	...		
...	...	...	...	...	...		
325929	0.00	C1322395030	0.00	34345.29	...		
325930	0.00	M1758068394	0.00	0.00	...		
325931	0.00	C1955063491	84106.06	97605.52	...		
325932	0.00	C1407924212	0.00	236678.98	...		
325933	228554.72	C763824199	389833.54	478305.71	...		
...	...	...	...	...	...		
	isCustomerDest	hour	day	abs_balance_diff	abs_errorBalanceDest	\	
0	1	17	16	581341.77	5.813418e+05		
1	1	10	14	74969.02	1.455192e-11		
2	1	10	1	152359.15	1.000000e-02		
3	1	14	12	553238.82	1.000000e-02		
4	1	10	5	865248.38	7.391534e+05		

```

...
325929      1      18      11      13965.29      0.000000e+00
325930      0      19      15      7962.77      8.234770e+03
325931      1      19      9      13499.45      1.000000e-02
325932      1      9      15      226660.98      0.000000e+00
325933      1      11      16      0.00      2.051711e+04

      is_empty_orig  is_empty_dest  amount_ratio_org  log_amount  red_flag
0              0              0      0.009495      12.579950      1
1              1              0      74969.020000      11.224844      1
2              1              0      152359.150000      11.934002      1
3              1              0      553238.820000      13.223547      1
4              0              0      22.876854      12.977627      1
...
325929      0              1      1.685162      10.444249      1
325930      0              1      30.163993      9.016242      1
325931      1              0      13499.450000      9.510478      1
325932      0              1      23.623014      12.374464      1
325933      0              0      0.322888      11.599014      1

```

[325934 rows x 25 columns]

```
[3]: # Kiểm tra cột
df.columns
```

```
[3]: Index(['Unnamed: 0', 'step', 'type', 'amount', 'nameOrig', 'oldbalanceOrg',
        'newbalanceOrig', 'nameDest', 'oldbalanceDest', 'newbalanceDest',
        'isFraud', 'isFlaggedFraud', 'balance_diff', 'errorBalanceDest',
        'isMerchant', 'isCustomerDest', 'hour', 'day', 'abs_balance_diff',
        'abs_errorBalanceDest', 'is_empty_orig', 'is_empty_dest',
        'amount_ratio_org', 'log_amount', 'red_flag'],
        dtype='object')
```

Feature	Leakage?	Loại leakage	Ghi chú
oldbalanceOrg / newbalanceOrig		Temporal + system	Drop
oldbalanceDest / newbalanceDest		Severe leakage	Drop
isFlaggedFraud		Label proxy	Drop
balance_diff		Accounting logic	Drop
errorBalanceDest		Hậu kiểm	Drop
abs_*		Inherited leakage	Drop
amount_ratio_org		Model-dependent	Drop for RF
is_empty_orig		Shortcut risk	RF danger
is_empty_dest		System artifact	Drop
red_flag		Composite proxy	Drop
hour / day		Safe	Keep
log_amount		Safe	Keep
type		Causal	Keep

```
[4]: # Drop các cột không cần thiết (giống như trong Random Forest)
drop_cols = [
    "Unnamed: 0",    # index giá
    "nameOrig",      # ID
    "nameDest",      # ID
    "isFlaggedFraud", # proxy của nhãn
    'oldbalanceOrg',
    'newbalanceOrig',
    'oldbalanceDest',
    'newbalanceDest',
    'step',
    'balance_diff',
    'errorBalanceDest', # NGUỒN LEAKAGE CHÍNH
    'abs_balance_diff',
    'abs_errorBalanceDest',
    'red_flag', # Derived from leakage
    'amount_ratio_org', # Tín hiệu rút sạch tiền quá lộ liễu
    'is_empty_orig', # Trạng thái tài khoản ban đầu
    'is_empty_dest' # Trạng thái tài khoản nhận
]

df = df.drop(columns=[c for c in drop_cols if c in df.columns])
print(f"\nShape sau khi drop: {df.shape}")
df.columns
```

Shape sau khi drop: (325934, 8)

```
[4]: Index(['type', 'amount', 'isFraud', 'isMerchant', 'isCustomerDest', 'hour',
          'day', 'log_amount'],
          dtype='object')
```

4 PHÂN TÍCH DỮ LIỆU

```
[5]: # Kiểm tra phân phối lớp
print("\nPhân phối lớp:")
print(df['isFraud'].value_counts())
print(f"Tỷ lệ fraud: {df['isFraud'].mean():.4f}")
```

Phân phối lớp:

isFraud

0 317721

1 8213

Name: count, dtype: int64

Tỷ lệ fraud: 0.0252

```
[6]: # Chuẩn bị features và target
target = "isFraud"
X = df.drop(columns=target)
y = df[target]

# Phân chia train-test
X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,
    random_state=42,
    stratify=y
)

print(f"\nKích thước tập train: {X_train.shape}")
print(f"Kích thước tập test: {X_test.shape}")
```

Kích thước tập train: (260747, 7)

Kích thước tập test: (65187, 7)

5 TIỀN XỬ LÝ FEATURES

```
[7]: # Phân loại feature
categorical_features = X.select_dtypes(include=['object', 'category']).columns.
    ↪tolist()
numeric_features = X.select_dtypes(include=['int64', 'float64']).columns.
    ↪tolist()

print(f"\nCategorical features: {categorical_features}")
print(f"Numeric features: {numeric_features}")

# Tạo preprocessor với OneHotEncoding cho categorical và StandardScaler cho
    ↪numeric
preprocessor = ColumnTransformer(
    transformers=[
        ("cat", OneHotEncoder(handle_unknown="ignore", drop="first"),
        ↪categorical_features),
        ("num", StandardScaler(), numeric_features) # XGBoost thường không cần
        ↪scaling nhưng có thể hữu ích
    ],
    remainder="drop"
)
```

Categorical features: ['type']

Numeric features: ['amount', 'isMerchant', 'isCustomerDest', 'hour', 'day',
'log_amount']

6 MÔ HÌNH XGBOOST CƠ BẢN

```
[8]: # Tính scale_pos_weight cho mất cân bằng
scale_pos_weight = len(y_train[y_train == 0]) / len(y_train[y_train == 1])
print(f"scale_pos_weight: {scale_pos_weight:.2f}")

# Pipeline cơ bản
xgb_baseline = Pipeline(steps=[
    ("preprocess", preprocessor),
    ("clf", xgb.XGBClassifier(
        n_estimators=300,
        max_depth=6,
        learning_rate=0.1,
        scale_pos_weight=scale_pos_weight,
        random_state=42,
        n_jobs=-1,
        eval_metric='logloss'
    ))
])
```

scale_pos_weight: 38.69

```
[9]: # Huấn luyện
xgb_baseline.fit(X_train, y_train)
```

```
[9]: Pipeline(steps=[('preprocess',
                      ColumnTransformer(transformers=[('cat',
                                                       OneHotEncoder(drop='first',
                                                                       handle_unknown='ignore'),
                                                       ['type']),
                                                       ('num', StandardScaler(),
                                                       ['amount', 'isMerchant',
                                                       'isCustomerDest', 'hour',
                                                       'day', 'log_amount'])])),
                      ('clf',
                       XGBClassifier(base_score=None, booster=None, callbacks=None,
                                      colsample_bylevel=None, colsample_bynode=None,
                                      colsample_bytree=None, feature_types=None, feature_weights=None,
                                      gamma=None, grow_policy=None, importance_type=None,
                                      interaction_constraints=None, learning_rate=0.1,
                                      max_bin=None, max_cat_threshold=None,
                                      max_cat_to_onehot=None, max_delta_step=None,
                                      max_depth=6, max_leaves=None,
                                      min_child_weight=None, missing=nan,
                                      monotone_constraints=None, multi_strategy=None,
                                      n_estimators=300, n_jobs=-1,
```

```
num_parallel_tree=None, ...)))
```

```
[10]: # Dự đoán
y_pred_baseline = xgb_baseline.predict(X_test)
y_prob_baseline = xgb_baseline.predict_proba(X_test)[: , 1]
```

```
[11]: # Đánh giá
print("\nĐánh giá mô hình XGBoost cơ bản:")
print("ROC-AUC:", roc_auc_score(y_test, y_prob_baseline))
print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred_baseline))
print("\nClassification Report:")
print(classification_report(y_test, y_pred_baseline, digits=4))
```

Đánh giá mô hình XGBoost cơ bản:
ROC-AUC: 0.9750635835486086

Confusion Matrix:
[[59916 3628]
 [211 1432]]

Classification Report:

	precision	recall	f1-score	support
0	0.9965	0.9429	0.9690	63544
1	0.2830	0.8716	0.4273	1643
accuracy			0.9411	65187
macro avg	0.6397	0.9072	0.6981	65187
weighted avg	0.9785	0.9411	0.9553	65187

7 TỐI ƯU SIÊU THAM SỐ VỚI OPTUNA

```
[12]: cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

def objective_xgb(trial):
    """Hàm objective cho Optuna với XGBoost"""
    xgb_params = {
        "n_estimators": trial.suggest_int("n_estimators", 100, 500),
        "max_depth": trial.suggest_int("max_depth", 3, 10),
        "learning_rate": trial.suggest_float("learning_rate", 0.01, 0.3,
        log=True),
        "subsample": trial.suggest_float("subsample", 0.6, 1.0),
        "colsample_bytree": trial.suggest_float("colsample_bytree", 0.6, 1.0),
        "min_child_weight": trial.suggest_int("min_child_weight", 1, 10),
```

```

        "gamma": trial.suggest_float("gamma", 0, 5),
        "reg_alpha": trial.suggest_float("reg_alpha", 0, 10),
        "reg_lambda": trial.suggest_float("reg_lambda", 0, 10),
        "scale_pos_weight": scale_pos_weight,
        "random_state": 42,
        "n_jobs": -1,
        "eval_metric": 'logloss'
    }

    xgb_model = xgb.XGBClassifier(**xgb_params)

    pipe = Pipeline(steps=[
        ("preprocess", preprocessor),
        ("clf", xgb_model)
    ])

    scores = cross_val_score(
        pipe,
        X_train, y_train,
        cv=cv,
        scoring="roc_auc",
        n_jobs=-1
    )

    return scores.mean()

```

```

[13]: # Chay Optuna
study_xgb = optuna.create_study(direction="maximize")
study_xgb.optimize(objective_xgb, n_trials=30, show_progress_bar=True)

```

[I 2025-12-30 14:36:14,052] A new study created in memory with name: no-name-20785c19-c828-45ca-a99d-8617dc422755

0%| | 0/30 [00:00<?, ?it/s]

[I 2025-12-30 14:36:29,888] Trial 0 finished with value: 0.972898176801402 and parameters: {'n_estimators': 250, 'max_depth': 4, 'learning_rate': 0.09984377239103541, 'subsample': 0.9907418495974626, 'colsample_bytree': 0.6396368621111381, 'min_child_weight': 1, 'gamma': 4.910486172747291, 'reg_alpha': 6.48508740142799, 'reg_lambda': 4.475815239800053}. Best is trial 0 with value: 0.972898176801402.

[I 2025-12-30 14:36:47,262] Trial 1 finished with value: 0.9720554814350457 and parameters: {'n_estimators': 256, 'max_depth': 8, 'learning_rate': 0.01620238559380132, 'subsample': 0.7929718146247484, 'colsample_bytree': 0.8566719112024781, 'min_child_weight': 10, 'gamma': 1.5425769115306287, 'reg_alpha': 4.404539218708332, 'reg_lambda': 6.694982295203765}. Best is trial 0 with value: 0.972898176801402.

[I 2025-12-30 14:37:08,112] Trial 2 finished with value: 0.9728135855391361 and parameters: {'n_estimators': 326, 'max_depth': 8, 'learning_rate':

0.018068757061556605, 'subsample': 0.8181955061534623, 'colsample_bytree': 0.9663256340322345, 'min_child_weight': 7, 'gamma': 3.193924739212286, 'reg_alpha': 9.312251047648205, 'reg_lambda': 0.942281155437733}. Best is trial 0 with value: 0.972898176801402.

[I 2025-12-30 14:37:21,861] Trial 3 finished with value: 0.9726617697951271 and parameters: {'n_estimators': 150, 'max_depth': 9, 'learning_rate': 0.0760260436667773, 'subsample': 0.6378589565164522, 'colsample_bytree': 0.7913024059872366, 'min_child_weight': 6, 'gamma': 2.56193363010944, 'reg_alpha': 1.841970347417008, 'reg_lambda': 8.79346328552128}. Best is trial 0 with value: 0.972898176801402.

[I 2025-12-30 14:37:43,292] Trial 4 finished with value: 0.976105110209296 and parameters: {'n_estimators': 432, 'max_depth': 8, 'learning_rate': 0.1318355672131095, 'subsample': 0.9386044842156706, 'colsample_bytree': 0.8655723536737898, 'min_child_weight': 1, 'gamma': 2.0985905368638047, 'reg_alpha': 7.227072603651484, 'reg_lambda': 2.47821053440322}. Best is trial 4 with value: 0.976105110209296.

[I 2025-12-30 14:37:52,797] Trial 5 finished with value: 0.9724199740996606 and parameters: {'n_estimators': 202, 'max_depth': 3, 'learning_rate': 0.2139530920675537, 'subsample': 0.8148760151205263, 'colsample_bytree': 0.6439230529428365, 'min_child_weight': 3, 'gamma': 4.07134662611138, 'reg_alpha': 6.612386498266001, 'reg_lambda': 2.982307912409543}. Best is trial 4 with value: 0.976105110209296.

[I 2025-12-30 14:38:17,488] Trial 6 finished with value: 0.9744052267463873 and parameters: {'n_estimators': 396, 'max_depth': 9, 'learning_rate': 0.12565132527584702, 'subsample': 0.8698844690377797, 'colsample_bytree': 0.6537565570298248, 'min_child_weight': 10, 'gamma': 0.7923631837361844, 'reg_alpha': 9.920735010243131, 'reg_lambda': 4.355073285922412}. Best is trial 4 with value: 0.976105110209296.

[I 2025-12-30 14:38:37,744] Trial 7 finished with value: 0.9747117598233974 and parameters: {'n_estimators': 327, 'max_depth': 9, 'learning_rate': 0.05833886609374307, 'subsample': 0.7875142488077472, 'colsample_bytree': 0.8095177160890217, 'min_child_weight': 5, 'gamma': 3.651230192149122, 'reg_alpha': 4.09770431819417, 'reg_lambda': 1.6557164830668147}. Best is trial 4 with value: 0.976105110209296.

[I 2025-12-30 14:39:04,171] Trial 8 finished with value: 0.9743016995095776 and parameters: {'n_estimators': 400, 'max_depth': 9, 'learning_rate': 0.1860859152716422, 'subsample': 0.8492543320334974, 'colsample_bytree': 0.7794490574992529, 'min_child_weight': 8, 'gamma': 0.41776172059682315, 'reg_alpha': 5.945421424625124, 'reg_lambda': 7.386553978446182}. Best is trial 4 with value: 0.976105110209296.

[I 2025-12-30 14:39:20,305] Trial 9 finished with value: 0.9750266296539299 and parameters: {'n_estimators': 277, 'max_depth': 10, 'learning_rate': 0.27336293289408325, 'subsample': 0.8817825083611126, 'colsample_bytree': 0.7780547175609492, 'min_child_weight': 9, 'gamma': 2.0880415134545425, 'reg_alpha': 5.29009209814652, 'reg_lambda': 6.6820935513438595}. Best is trial 4 with value: 0.976105110209296.

[I 2025-12-30 14:39:44,222] Trial 10 finished with value: 0.9745005960588854 and parameters: {'n_estimators': 500, 'max_depth': 6, 'learning_rate':

0.031661378737547374, 'subsample': 0.9366097941008074, 'colsample_bytree': 0.9557789173172895, 'min_child_weight': 1, 'gamma': 1.563347021279693, 'reg_alpha': 7.791761513270265, 'reg_lambda': 2.6367703894092958}. Best is trial 4 with value: 0.976105110209296.

[I 2025-12-30 14:40:05,853] Trial 11 finished with value: 0.9755315622062335 and parameters: {'n_estimators': 499, 'max_depth': 10, 'learning_rate': 0.2836207228911647, 'subsample': 0.9331219864570522, 'colsample_bytree': 0.8793737656037628, 'min_child_weight': 4, 'gamma': 2.014121208511816, 'reg_alpha': 2.8444085917235826, 'reg_lambda': 6.064552170624216}. Best is trial 4 with value: 0.976105110209296.

[I 2025-12-30 14:40:26,328] Trial 12 finished with value: 0.9750123580808644 and parameters: {'n_estimators': 490, 'max_depth': 6, 'learning_rate': 0.13930499044894817, 'subsample': 0.974152254874804, 'colsample_bytree': 0.8969185000422475, 'min_child_weight': 3, 'gamma': 2.774108208687745, 'reg_alpha': 1.991472320744481, 'reg_lambda': 4.908140045561549}. Best is trial 4 with value: 0.976105110209296.

[I 2025-12-30 14:40:54,216] Trial 13 finished with value: 0.9738694045705778 and parameters: {'n_estimators': 430, 'max_depth': 7, 'learning_rate': 0.27629535257916554, 'subsample': 0.7220206937568433, 'colsample_bytree': 0.8935089350686064, 'min_child_weight': 3, 'gamma': 1.6915565506662298, 'reg_alpha': 3.109655802984846, 'reg_lambda': 0.417867057614413}. Best is trial 4 with value: 0.976105110209296.

[I 2025-12-30 14:41:27,978] Trial 14 finished with value: 0.9730680634815794 and parameters: {'n_estimators': 443, 'max_depth': 10, 'learning_rate': 0.03874661489299247, 'subsample': 0.9177283897857726, 'colsample_bytree': 0.7192054757985531, 'min_child_weight': 4, 'gamma': 1.0838215787798693, 'reg_alpha': 0.35047388003326585, 'reg_lambda': 3.2456646101132005}. Best is trial 4 with value: 0.976105110209296.

[I 2025-12-30 14:41:51,748] Trial 15 finished with value: 0.9753185655224635 and parameters: {'n_estimators': 365, 'max_depth': 7, 'learning_rate': 0.1638789403905161, 'subsample': 0.9325434585509994, 'colsample_bytree': 0.8634024941654763, 'min_child_weight': 2, 'gamma': 0.09537970550156682, 'reg_alpha': 8.06135530041619, 'reg_lambda': 9.545029477176318}. Best is trial 4 with value: 0.976105110209296.

[I 2025-12-30 14:42:15,542] Trial 16 finished with value: 0.9741604652011361 and parameters: {'n_estimators': 433, 'max_depth': 5, 'learning_rate': 0.08917385503123548, 'subsample': 0.7085538018489852, 'colsample_bytree': 0.9995322697603539, 'min_child_weight': 5, 'gamma': 2.1214980872020357, 'reg_alpha': 3.1862028162779708, 'reg_lambda': 5.8564736848071135}. Best is trial 4 with value: 0.976105110209296.

[I 2025-12-30 14:42:34,964] Trial 17 finished with value: 0.976238899363756 and parameters: {'n_estimators': 464, 'max_depth': 8, 'learning_rate': 0.292031438223591, 'subsample': 0.9614681618171416, 'colsample_bytree': 0.9328331701255683, 'min_child_weight': 2, 'gamma': 3.10450780457985, 'reg_alpha': 7.625850948580585, 'reg_lambda': 7.797483967720321}. Best is trial 17 with value: 0.976238899363756.

[I 2025-12-30 14:42:54,189] Trial 18 finished with value: 0.9746561884215689 and parameters: {'n_estimators': 369, 'max_depth': 8, 'learning_rate':

0.05926564065297405, 'subsample': 0.9819246581869342, 'colsample_bytree': 0.92948335802007, 'min_child_weight': 1, 'gamma': 3.3052411987746835, 'reg_alpha': 7.823904851828373, 'reg_lambda': 8.556771626710168}. Best is trial 17 with value: 0.976238899363756.

[I 2025-12-30 14:43:18,623] Trial 19 finished with value: 0.9758074197884247 and parameters: {'n_estimators': 463, 'max_depth': 7, 'learning_rate': 0.12976695930854812, 'subsample': 0.8915088206885384, 'colsample_bytree': 0.8309646300130742, 'min_child_weight': 2, 'gamma': 4.137678696758826, 'reg_alpha': 8.85502581096902, 'reg_lambda': 1.9009590818855486}. Best is trial 17 with value: 0.976238899363756.

[I 2025-12-30 14:43:26,174] Trial 20 finished with value: 0.9740795052788342 and parameters: {'n_estimators': 120, 'max_depth': 6, 'learning_rate': 0.18755910759583713, 'subsample': 0.9590473483048843, 'colsample_bytree': 0.7091500788038106, 'min_child_weight': 2, 'gamma': 2.933242335788117, 'reg_alpha': 6.931394430625345, 'reg_lambda': 7.661723678006316}. Best is trial 17 with value: 0.976238899363756.

[I 2025-12-30 14:43:50,483] Trial 21 finished with value: 0.9757441463702594 and parameters: {'n_estimators': 460, 'max_depth': 7, 'learning_rate': 0.11430694944400334, 'subsample': 0.8988427305325903, 'colsample_bytree': 0.8487319216014654, 'min_child_weight': 2, 'gamma': 4.346326867994611, 'reg_alpha': 8.845713700594011, 'reg_lambda': 1.918398219392468}. Best is trial 17 with value: 0.976238899363756.

[I 2025-12-30 14:44:16,083] Trial 22 finished with value: 0.9723005254133884 and parameters: {'n_estimators': 457, 'max_depth': 8, 'learning_rate': 0.010780354927962027, 'subsample': 0.8543642366394156, 'colsample_bytree': 0.8253777987971389, 'min_child_weight': 2, 'gamma': 3.8219126025232404, 'reg_alpha': 8.422552349364688, 'reg_lambda': 3.8243540073777105}. Best is trial 17 with value: 0.976238899363756.

[I 2025-12-30 14:44:38,291] Trial 23 finished with value: 0.9758062093633167 and parameters: {'n_estimators': 405, 'max_depth': 7, 'learning_rate': 0.0744330124355066, 'subsample': 0.902214278084932, 'colsample_bytree': 0.926819541894863, 'min_child_weight': 1, 'gamma': 4.528810506586791, 'reg_alpha': 7.369229190214933, 'reg_lambda': 1.7861059459931852}. Best is trial 17 with value: 0.976238899363756.

[I 2025-12-30 14:44:59,606] Trial 24 finished with value: 0.9745266182611809 and parameters: {'n_estimators': 465, 'max_depth': 5, 'learning_rate': 0.13861933982445176, 'subsample': 0.9543352239764129, 'colsample_bytree': 0.7447053698112812, 'min_child_weight': 4, 'gamma': 3.386976603006071, 'reg_alpha': 9.973278725797938, 'reg_lambda': 0.18595766392700686}. Best is trial 17 with value: 0.976238899363756.

[I 2025-12-30 14:45:19,652] Trial 25 finished with value: 0.9737164410656473 and parameters: {'n_estimators': 364, 'max_depth': 8, 'learning_rate': 0.03648886077420128, 'subsample': 0.9995081454279364, 'colsample_bytree': 0.9231621520170811, 'min_child_weight': 3, 'gamma': 2.5757629578123433, 'reg_alpha': 5.707322446363093, 'reg_lambda': 2.4497491160396887}. Best is trial 17 with value: 0.976238899363756.

[I 2025-12-30 14:45:40,021] Trial 26 finished with value: 0.9743864248292509 and parameters: {'n_estimators': 417, 'max_depth': 5, 'learning_rate':

```
0.20550658773598665, 'subsample': 0.901633093473473, 'colsample_bytree':
0.8235424463505985, 'min_child_weight': 2, 'gamma': 4.980261471798712,
'reg_alpha': 8.982071921207858, 'reg_lambda': 3.729050198241426}. Best is trial
17 with value: 0.976238899363756.
[I 2025-12-30 14:46:09,767] Trial 27 finished with value: 0.9754460036912324 and
parameters: {'n_estimators': 472, 'max_depth': 7, 'learning_rate':
0.2388112088005547, 'subsample': 0.7552758594578974, 'colsample_bytree':
0.9741635369199463, 'min_child_weight': 1, 'gamma': 3.8093226752452347,
'reg_alpha': 7.206528088906138, 'reg_lambda': 1.0733242263729401}. Best is trial
17 with value: 0.976238899363756.
[I 2025-12-30 14:46:32,019] Trial 28 finished with value: 0.9756296853274096 and
parameters: {'n_estimators': 332, 'max_depth': 8, 'learning_rate':
0.15198672584263864, 'subsample': 0.8385314256664366, 'colsample_bytree':
0.9123352177615902, 'min_child_weight': 4, 'gamma': 3.0102010720356955,
'reg_alpha': 8.344742326911003, 'reg_lambda': 5.723148182920532}. Best is trial
17 with value: 0.976238899363756.
[I 2025-12-30 14:46:49,406] Trial 29 finished with value: 0.9728490415942149 and
parameters: {'n_estimators': 372, 'max_depth': 4, 'learning_rate':
0.04841556167679168, 'subsample': 0.9725501279442598, 'colsample_bytree':
0.836254977928879, 'min_child_weight': 1, 'gamma': 4.6441231134946115,
'reg_alpha': 6.386543143217574, 'reg_lambda': 4.906227310572705}. Best is trial
17 with value: 0.976238899363756.
```

```
[14]: # Kết quả tối ưu XGBoost
print("Best ROC-AUC (CV):", study_xgb.best_value)
print("Best params:", study_xgb.best_params)
```

```
Best ROC-AUC (CV): 0.976238899363756
Best params: {'n_estimators': 464, 'max_depth': 8, 'learning_rate':
0.292031438223591, 'subsample': 0.9614681618171416, 'colsample_bytree':
0.9328331701255683, 'min_child_weight': 2, 'gamma': 3.10450780457985,
'reg_alpha': 7.625850948580585, 'reg_lambda': 7.797483967720321}
```

8 MÔ HÌNH XGBOOST VỚI THAM SỐ TỐI ƯU

```
[15]: best_xgb_params = study_xgb.best_params.copy()
# Thêm các tham số cố định
best_xgb_params.update({
    'scale_pos_weight': scale_pos_weight,
    'random_state': 42,
    'n_jobs': -1,
    'eval_metric': 'logloss'
})

best_xgb = xgb.XGBClassifier(**best_xgb_params)
best_xgb_model = Pipeline(steps=[
    ("preprocess", preprocessor),
```

```
    ("clf", best_xgb)
])
```

```
[16]: best_xgb_model.fit(X_train, y_train)
```

```
[16]: Pipeline(steps=[('preprocess',
                        ColumnTransformer(transformers=[('cat',
                                                         OneHotEncoder(drop='first',
                                                         handle_unknown='ignore'),
                                                         ['type']),
                                                         ('num', StandardScaler(),
                                                         ['amount', 'isMerchant',
                                                         'isCustomerDest', 'hour',
                                                         'day', 'log_amount'])])),
                      ('clf',
                       XGBClassifier(base_score=None, booster=None, callbacks=None,
                                     colsample_bylevel=None, colsample_bynode=None,
                                     colsample_bytr...
                                     gamma=3.10450780457985, grow_policy=None,
                                     importance_type=None,
                                     interaction_constraints=None,
                                     learning_rate=0.292031438223591, max_bin=None,
                                     max_cat_threshold=None, max_cat_to_onehot=None,
                                     max_delta_step=None, max_depth=8,
                                     max_leaves=None, min_child_weight=2, missing=nan,
                                     monotone_constraints=None, multi_strategy=None,
                                     n_estimators=464, n_jobs=-1,
                                     num_parallel_tree=None, ...))])
```

```
[17]: # Đánh giá mô hình tối ưu
y_pred_best_xgb = best_xgb_model.predict(X_test)
y_prob_best_xgb = best_xgb_model.predict_proba(X_test)[:, 1]

print("\nĐánh giá XGBoost tối ưu:")
print("Test ROC-AUC:", roc_auc_score(y_test, y_prob_best_xgb))
print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred_best_xgb))
print("\nClassification Report:")
print(classification_report(y_test, y_pred_best_xgb, digits=4))
```

Đánh giá XGBoost tối ưu:

Test ROC-AUC: 0.9786812885233951

Confusion Matrix:

```
[[60809  2735]
 [  228 1415]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.9963	0.9570	0.9762	63544
1	0.3410	0.8612	0.4885	1643
accuracy			0.9545	65187
macro avg	0.6686	0.9091	0.7324	65187
weighted avg	0.9797	0.9545	0.9639	65187

```
[18]: # Feature importance
feature_names = best_xgb_model.named_steps["preprocess"].get_feature_names_out()
importances = best_xgb_model.named_steps["clf"].feature_importances_

fi_xgb = pd.DataFrame({
    "feature": feature_names,
    "importance": importances
}).sort_values(by="importance", ascending=False)

print("\nTop 10 Feature Importance (XGBoost):")
print(fi_xgb.head(10))
```

Top 10 Feature Importance (XGBoost):

	feature	importance
2	cat__type_PAYMENT	0.399967
3	cat__type_TRANSFER	0.258792
0	cat__type_CASH_OUT	0.153248
5	num__isMerchant	0.130093
7	num__hour	0.015112
4	num__amount	0.012985
9	num__log_amount	0.011780
8	num__day	0.010872
1	cat__type_DEBIT	0.007151
6	num__isCustomerDest	0.000000

9 THỬ NGHIỆM RESAMPLING STRATEGIES

```
[19]: strategies = [
    ('smote', SMOTE(random_state=42)),
    ('smote_enn', SMOTEENN(random_state=42)),
    ('smote_tomek', SMOTETomek(random_state=42)),
    ('oversample', RandomOverSampler(random_state=42)),
    ('undersample', RandomUnderSampler(random_state=42)),
]
```

```

results_xgb = {}

for strategy_name, sampler in strategies:
    print(f"\nĐang thử: {strategy_name.upper()}")

    # Tạo pipeline với resampling (không dùng scale_pos_weight khi đã
    ↪resampling)
    pipe = ImbPipeline(steps=[
        ("preprocess", preprocessor),
        ("resample", sampler),
        ("clf", xgb.XGBClassifier(
            n_estimators=100, # Dùng ít trees để test nhanh
            max_depth=6,
            learning_rate=0.1,
            scale_pos_weight=1, # Không dùng vì đã resampling
            random_state=42,
            n_jobs=-1,
            eval_metric='logloss'
        ))
    ])

    # Huấn luyện
    pipe.fit(X_train, y_train)

    # Dự đoán
    y_pred = pipe.predict(X_test)
    y_prob = pipe.predict_proba(X_test)[:, 1]

    # Tính metrics
    roc_auc = roc_auc_score(y_test, y_prob)
    pr_auc = average_precision_score(y_test, y_prob)
    recall = recall_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)

    # Lưu kết quả
    results_xgb[strategy_name] = {
        'roc_auc': roc_auc,
        'pr_auc': pr_auc,
        'recall': recall,
        'precision': precision,
        'f1': f1,
        'model': pipe
    }

    print(f" ROC-AUC: {roc_auc:.4f}")
    print(f" PR-AUC: {pr_auc:.4f}")

```

```
print(f" Recall: {recall:.4f}")
print(f" Precision: {precision:.4f}")
print(f" F1-Score: {f1:.4f}")
```

Đang thử: SMOTE

```
ROC-AUC: 0.9710
PR-AUC: 0.7539
Recall: 0.8217
Precision: 0.3496
F1-Score: 0.4905
```

Đang thử: SMOTE_ENN

```
ROC-AUC: 0.9710
PR-AUC: 0.7299
Recall: 0.8332
Precision: 0.3280
F1-Score: 0.4707
```

Đang thử: SMOTE_TOMEK

```
ROC-AUC: 0.9712
PR-AUC: 0.7547
Recall: 0.8162
Precision: 0.3491
F1-Score: 0.4891
```

Đang thử: OVERSAMPLE

```
ROC-AUC: 0.9715
PR-AUC: 0.7623
Recall: 0.8801
Precision: 0.2240
F1-Score: 0.3571
```

Đang thử: UNDERSAMPLE

```
ROC-AUC: 0.9711
PR-AUC: 0.7387
Recall: 0.8898
Precision: 0.2158
F1-Score: 0.3474
```

```
[20]: # Chọn chiến lược tốt nhất dựa trên F1-score
best_strategy_xgb = max(results_xgb.keys(), key=lambda x: results_xgb[x]['f1'])
print(f"\nChiến lược resampling tốt nhất cho XGBoost: {best_strategy_xgb.
    ↪upper()} (F1 = {results_xgb[best_strategy_xgb]['f1']:.4f})")
```

Chiến lược resampling tốt nhất cho XGBoost: SMOTE (F1 = 0.4905)

10 TỐI ƯU XGBOOST VỚI RESAMPLING TỐT NHẤT

```
[21]: print(f"TỐI ƯU XGBOOST VỚI {best_strategy_xgb.upper()}")

# Lấy sampler tốt nhất
best_sampler_xgb = strategies[[s[0] for s in strategies].
    ↪index(best_strategy_xgb)][1]

def objective_xgb_resample(trial):
    """Hàm objective cho Optuna với XGBoost + Resampling"""
    # Hyperparameters cho XGBoost
    xgb_params = {
        "n_estimators": trial.suggest_int("n_estimators", 100, 500),
        "max_depth": trial.suggest_int("max_depth", 3, 10),
        "learning_rate": trial.suggest_float("learning_rate", 0.01, 0.3,
    ↪log=True),
        "subsample": trial.suggest_float("subsample", 0.6, 1.0),
        "colsample_bytree": trial.suggest_float("colsample_bytree", 0.6, 1.0),
        "min_child_weight": trial.suggest_int("min_child_weight", 1, 10),
        "gamma": trial.suggest_float("gamma", 0, 5),
        "reg_alpha": trial.suggest_float("reg_alpha", 0, 10),
        "reg_lambda": trial.suggest_float("reg_lambda", 0, 10),
        "random_state": 42,
        "n_jobs": -1,
        "eval_metric": 'logloss',
        "scale_pos_weight": 1 # Không dùng vì đã resampling
    }

    # Hyperparameters cho SMOTE (nếu dùng)
    if best_strategy_xgb == 'smote':
        k_neighbors = trial.suggest_int("k_neighbors", 3, 10)
        sampling_strategy = trial.suggest_float("sampling_strategy", 0.3, 0.9)
        sampler = SMOTE(
            random_state=42,
            k_neighbors=k_neighbors,
            sampling_strategy=sampling_strategy
        )
    else:
        sampler = best_sampler_xgb

    # Tạo pipeline
    pipe = ImbPipeline(steps=[
        ("preprocess", preprocessor),
        ("resample", sampler),
        ("clf", xgb.XGBClassifier(**xgb_params))
    ])
    ↪
```

```

# Cross-validation
scores = cross_val_score(
    pipe, X_train, y_train,
    cv=StratifiedKFold(n_splits=3, shuffle=True, random_state=42),
    scoring='roc_auc',
    n_jobs=-1
)
return scores.mean()

```

TỐI ƯU XGBOOST VỚI SMOTE

```

[22]: # Chạy Optuna
study_xgb_resample = optuna.create_study(direction="maximize")
study_xgb_resample.optimize(objective_xgb_resample, n_trials=20,
    ↪show_progress_bar=True)

```

[I 2025-12-30 14:50:32,723] A new study created in memory with name: no-name-2bd3c78e-eb9a-4d9c-8d7e-6b33a306eb60

0%| | 0/20 [00:00<?, ?it/s]

[I 2025-12-30 14:50:49,888] Trial 0 finished with value: 0.9713877497991015 and parameters: {'n_estimators': 279, 'max_depth': 10, 'learning_rate': 0.14934949943566778, 'subsample': 0.6465549434736566, 'colsample_bytree': 0.6859112761482669, 'min_child_weight': 3, 'gamma': 2.9413045400551523, 'reg_alpha': 0.04431350287012026, 'reg_lambda': 5.87597786653707, 'k_neighbors': 6, 'sampling_strategy': 0.5255416116453809}. Best is trial 0 with value: 0.9713877497991015.

[I 2025-12-30 14:51:14,089] Trial 1 finished with value: 0.9719039305847881 and parameters: {'n_estimators': 377, 'max_depth': 4, 'learning_rate': 0.13595739838205656, 'subsample': 0.97406529180431, 'colsample_bytree': 0.8474340663769374, 'min_child_weight': 2, 'gamma': 2.0255736377438804, 'reg_alpha': 0.5549837569757032, 'reg_lambda': 7.27485662728812, 'k_neighbors': 10, 'sampling_strategy': 0.7486939302853053}. Best is trial 1 with value: 0.9719039305847881.

[I 2025-12-30 14:51:34,764] Trial 2 finished with value: 0.9713345939680464 and parameters: {'n_estimators': 238, 'max_depth': 8, 'learning_rate': 0.03540341620282129, 'subsample': 0.9667927083476086, 'colsample_bytree': 0.7383512516615692, 'min_child_weight': 10, 'gamma': 4.272584002597652, 'reg_alpha': 9.452457081244582, 'reg_lambda': 9.975795107472482, 'k_neighbors': 10, 'sampling_strategy': 0.6794410345816426}. Best is trial 1 with value: 0.9719039305847881.

[I 2025-12-30 14:52:05,022] Trial 3 finished with value: 0.9713485743466048 and parameters: {'n_estimators': 299, 'max_depth': 10, 'learning_rate': 0.0350613066898389, 'subsample': 0.7701487802575331, 'colsample_bytree': 0.6278362056017794, 'min_child_weight': 6, 'gamma': 2.7005179248495197, 'reg_alpha': 3.6072585213933515, 'reg_lambda': 4.498342995174917, 'k_neighbors': 4, 'sampling_strategy': 0.7971172712458555}. Best is trial 1 with value: 0.9719039305847881.

[I 2025-12-30 14:52:35,260] Trial 4 finished with value: 0.9714074793812587 and parameters: {'n_estimators': 361, 'max_depth': 7, 'learning_rate': 0.021769773455837527, 'subsample': 0.9429515265327706, 'colsample_bytree': 0.6501315241884337, 'min_child_weight': 7, 'gamma': 4.717791469027731, 'reg_alpha': 9.00544665950188, 'reg_lambda': 5.27027807119723, 'k_neighbors': 5, 'sampling_strategy': 0.7918073373483014}. Best is trial 1 with value: 0.9719039305847881.

[I 2025-12-30 14:52:50,806] Trial 5 finished with value: 0.9717990262524766 and parameters: {'n_estimators': 201, 'max_depth': 4, 'learning_rate': 0.29648023603916246, 'subsample': 0.9364762017121436, 'colsample_bytree': 0.9853828595342103, 'min_child_weight': 8, 'gamma': 4.376112560284101, 'reg_alpha': 6.39678571686654, 'reg_lambda': 5.739672461932189, 'k_neighbors': 7, 'sampling_strategy': 0.5604087156329791}. Best is trial 1 with value: 0.9719039305847881.

[I 2025-12-30 14:53:25,018] Trial 6 finished with value: 0.970409277631311 and parameters: {'n_estimators': 379, 'max_depth': 10, 'learning_rate': 0.06828493528545743, 'subsample': 0.6140028159921086, 'colsample_bytree': 0.8290631445732046, 'min_child_weight': 9, 'gamma': 0.8042399946746509, 'reg_alpha': 3.7918010142623, 'reg_lambda': 1.1338229743408634, 'k_neighbors': 7, 'sampling_strategy': 0.37887488328781416}. Best is trial 1 with value: 0.9719039305847881.

[I 2025-12-30 14:53:55,396] Trial 7 finished with value: 0.9708882083895071 and parameters: {'n_estimators': 400, 'max_depth': 9, 'learning_rate': 0.07459716756094048, 'subsample': 0.9948040224347491, 'colsample_bytree': 0.9095927224803233, 'min_child_weight': 1, 'gamma': 1.9715698356695759, 'reg_alpha': 5.036812869330674, 'reg_lambda': 3.1032205149964307, 'k_neighbors': 4, 'sampling_strategy': 0.6352397466324697}. Best is trial 1 with value: 0.9719039305847881.

[I 2025-12-30 14:54:38,571] Trial 8 finished with value: 0.9705915311771327 and parameters: {'n_estimators': 384, 'max_depth': 10, 'learning_rate': 0.0323363247320671, 'subsample': 0.8949821587680664, 'colsample_bytree': 0.8995120738644491, 'min_child_weight': 4, 'gamma': 0.4553814821629759, 'reg_alpha': 5.719970518116079, 'reg_lambda': 9.69446300145285, 'k_neighbors': 5, 'sampling_strategy': 0.8739173664370725}. Best is trial 1 with value: 0.9719039305847881.

[I 2025-12-30 14:55:00,913] Trial 9 finished with value: 0.9722341709065869 and parameters: {'n_estimators': 329, 'max_depth': 7, 'learning_rate': 0.03060279743436547, 'subsample': 0.6621825127169095, 'colsample_bytree': 0.7489313703131134, 'min_child_weight': 2, 'gamma': 0.9879482604259149, 'reg_alpha': 1.11529378058285, 'reg_lambda': 5.98093372320315, 'k_neighbors': 3, 'sampling_strategy': 0.3895526603734732}. Best is trial 9 with value: 0.9722341709065869.

[I 2025-12-30 14:55:13,966] Trial 10 finished with value: 0.9671903115330137 and parameters: {'n_estimators': 103, 'max_depth': 6, 'learning_rate': 0.011009902932256496, 'subsample': 0.7323618179224383, 'colsample_bytree': 0.756908793566713, 'min_child_weight': 4, 'gamma': 0.013206328859299954, 'reg_alpha': 2.0614657196676944, 'reg_lambda': 7.742466348824221, 'k_neighbors': 3, 'sampling_strategy': 0.31153340273183205}. Best is trial 9 with value: 0.9722341709065869.

0.9722341709065869.

[I 2025-12-30 14:55:38,388] Trial 11 finished with value: 0.9714362502298711 and parameters: {'n_estimators': 487, 'max_depth': 3, 'learning_rate': 0.12048580720225457, 'subsample': 0.8546741697178838, 'colsample_bytree': 0.8205983776508292, 'min_child_weight': 1, 'gamma': 1.5140990901891207, 'reg_alpha': 0.3446719635213186, 'reg_lambda': 7.761037709468521, 'k_neighbors': 9, 'sampling_strategy': 0.42347562742990513}. Best is trial 9 with value: 0.9722341709065869.

[I 2025-12-30 14:56:09,096] Trial 12 finished with value: 0.9715093149442481 and parameters: {'n_estimators': 459, 'max_depth': 5, 'learning_rate': 0.2000274169446814, 'subsample': 0.6919247067780613, 'colsample_bytree': 0.8766760327384491, 'min_child_weight': 3, 'gamma': 1.3525488619734047, 'reg_alpha': 1.8897901496033882, 'reg_lambda': 7.8344300097751445, 'k_neighbors': 9, 'sampling_strategy': 0.7108374407384054}. Best is trial 9 with value: 0.9722341709065869.

[I 2025-12-30 14:56:30,699] Trial 13 finished with value: 0.9704862321594826 and parameters: {'n_estimators': 329, 'max_depth': 6, 'learning_rate': 0.016136766692619858, 'subsample': 0.8245028706288671, 'colsample_bytree': 0.7541209092936652, 'min_child_weight': 2, 'gamma': 3.272758881426338, 'reg_alpha': 1.7891730386444156, 'reg_lambda': 6.887260093373682, 'k_neighbors': 8, 'sampling_strategy': 0.47345648392190487}. Best is trial 9 with value: 0.9722341709065869.

[I 2025-12-30 14:56:56,978] Trial 14 finished with value: 0.9718461944058584 and parameters: {'n_estimators': 427, 'max_depth': 3, 'learning_rate': 0.09814843779399812, 'subsample': 0.7759538644001149, 'colsample_bytree': 0.84614460644326, 'min_child_weight': 5, 'gamma': 1.9497111732269394, 'reg_alpha': 3.5070198333294256, 'reg_lambda': 4.111398211601843, 'k_neighbors': 10, 'sampling_strategy': 0.750335550652067}. Best is trial 9 with value: 0.9722341709065869.

[I 2025-12-30 14:57:11,362] Trial 15 finished with value: 0.9720981343252957 and parameters: {'n_estimators': 197, 'max_depth': 5, 'learning_rate': 0.05717927496166805, 'subsample': 0.6906499593660664, 'colsample_bytree': 0.9877592060658691, 'min_child_weight': 2, 'gamma': 3.6352885977591285, 'reg_alpha': 1.076027721399176, 'reg_lambda': 2.791300595587792, 'k_neighbors': 3, 'sampling_strategy': 0.6069384216254068}. Best is trial 9 with value: 0.9722341709065869.

[I 2025-12-30 14:57:24,299] Trial 16 finished with value: 0.9715468804854739 and parameters: {'n_estimators': 155, 'max_depth': 7, 'learning_rate': 0.041646830033659384, 'subsample': 0.6840215094646273, 'colsample_bytree': 0.9968931357271774, 'min_child_weight': 5, 'gamma': 3.6383430846498532, 'reg_alpha': 7.508283387575401, 'reg_lambda': 1.5996528744041891, 'k_neighbors': 3, 'sampling_strategy': 0.6036454502328141}. Best is trial 9 with value: 0.9722341709065869.

[I 2025-12-30 14:57:38,081] Trial 17 finished with value: 0.9704098981812178 and parameters: {'n_estimators': 239, 'max_depth': 5, 'learning_rate': 0.02063332546149433, 'subsample': 0.6031989974235532, 'colsample_bytree': 0.9431056576260787, 'min_child_weight': 2, 'gamma': 3.675541879123415, 'reg_alpha': 2.8382927585466073, 'reg_lambda': 3.158697921231812, 'k_neighbors':

4, 'sampling_strategy': 0.3013590580442902}. Best is trial 9 with value: 0.9722341709065869.

[I 2025-12-30 14:57:53,023] Trial 18 finished with value: 0.9719021609426942 and parameters: {'n_estimators': 185, 'max_depth': 8, 'learning_rate': 0.053104440702238784, 'subsample': 0.7123817239894754, 'colsample_bytree': 0.702324083852369, 'min_child_weight': 3, 'gamma': 1.0527847423148167, 'reg_alpha': 1.0838375929623811, 'reg_lambda': 0.06967818477404641, 'k_neighbors': 3, 'sampling_strategy': 0.5066296611422421}. Best is trial 9 with value: 0.9722341709065869.

[I 2025-12-30 14:58:02,345] Trial 19 finished with value: 0.9674613811375847 and parameters: {'n_estimators': 129, 'max_depth': 5, 'learning_rate': 0.026616157815783095, 'subsample': 0.6421229867572578, 'colsample_bytree': 0.7878829690412824, 'min_child_weight': 1, 'gamma': 2.4357254067546394, 'reg_alpha': 2.7120317991770992, 'reg_lambda': 2.845070273358413, 'k_neighbors': 6, 'sampling_strategy': 0.4277152120674152}. Best is trial 9 with value: 0.9722341709065869.

```
[23]: print(f"\nBest ROC-AUC (CV): {study_xgb_resample.best_value:.6f}")
      print(f"Best parameters: {study_xgb_resample.best_params}")
```

Best ROC-AUC (CV): 0.972234
 Best parameters: {'n_estimators': 329, 'max_depth': 7, 'learning_rate': 0.03060279743436547, 'subsample': 0.6621825127169095, 'colsample_bytree': 0.7489313703131134, 'min_child_weight': 2, 'gamma': 0.9879482604259149, 'reg_alpha': 1.11529378058285, 'reg_lambda': 5.98093372320315, 'k_neighbors': 3, 'sampling_strategy': 0.3895526603734732}

11 HUẤN LUYỆN MÔ HÌNH CUỐI CÙNG

```
[24]: # Lấy best parameters
      best_params_xgb = study_xgb_resample.best_params

      # Tạo sampler với best parameters (nếu có)
      if best_strategy_xgb == 'smote':
          sampler_final = SMOTE(
              random_state=42,
              k_neighbors=best_params_xgb.get('k_neighbors', 5),
              sampling_strategy=best_params_xgb.get('sampling_strategy', 0.5)
          )
      else:
          sampler_final = best_sampler_xgb
```

```
[25]: # Tạo XGBoost với best parameters
      xgb_final_params = {
          'n_estimators': best_params_xgb.get('n_estimators', 300),
          'max_depth': best_params_xgb.get('max_depth', 6),
```

```

'learning_rate': best_params_xgb.get('learning_rate', 0.1),
'subsample': best_params_xgb.get('subsample', 1.0),
'colsample_bytree': best_params_xgb.get('colsample_bytree', 1.0),
'min_child_weight': best_params_xgb.get('min_child_weight', 1),
'gamma': best_params_xgb.get('gamma', 0),
'reg_alpha': best_params_xgb.get('reg_alpha', 0),
'reg_lambda': best_params_xgb.get('reg_lambda', 1),
'random_state': 42,
'n_jobs': -1,
'eval_metric': 'logloss',
'scale_pos_weight': 1
}

```

```

[26]: # Tạo final pipeline
final_pipe_xgb = ImbPipeline(steps=[
    ("preprocess", preprocessor),
    ("resample", sampler_final),
    ("clf", xgb.XGBClassifier(**xgb_final_params))
])

```

```

[27]: # Huấn luyện
final_pipe_xgb.fit(X_train, y_train)

```

```

[27]: Pipeline(steps=[('preprocess',
                        ColumnTransformer(transformers=[('cat',
                                                         OneHotEncoder(drop='first',
                                                         handle_unknown='ignore'),
                                                         ['type']),
                                                         ('num', StandardScaler(),
                                                         ['amount', 'isMerchant',
                                                         'isCustomerDest', 'hour',
                                                         'day', 'log_amount'])])),
                      ('resample',
                       SMOTE(k_neighbors=3, random_state=42,
                              sampling_strategy=0.3895526603734732)),
                      ('clf',
                       XGBClassifier(base_score=None,...
                                     gamma=0.9879482604259149, grow_policy=None,
                                     importance_type=None,
                                     interaction_constraints=None,
                                     learning_rate=0.03060279743436547, max_bin=None,
                                     max_cat_threshold=None, max_cat_to_onehot=None,
                                     max_delta_step=None, max_depth=7,
                                     max_leaves=None, min_child_weight=2, missing=nan,
                                     monotone_constraints=None, multi_strategy=None,
                                     n_estimators=329, n_jobs=-1,
                                     num_parallel_tree=None, ...))])

```

11.0.1 Chiến lược chọn điểm cắt (Threshold Tuning)

Với mô hình xác suất, ngưỡng mặc định 0.5 hiếm khi là tối ưu cho bài toán Fraud. Chúng ta phân tích hai kịch bản:

1. **Kịch bản 1: Maximize F1-Score (Cân bằng)**
 - Tìm điểm cắt để hài hòa giữa Precision và Recall.
 - Mục tiêu: Mô hình tổng quát tốt nhất.
2. **Kịch bản 2: High Recall (Ưu tiên bắt sót)**
 - Đặt mục tiêu Recall ≥ 0.9 (bắt được 90% gian lận).
 - Chấp nhận Precision thấp hơn (nhiều báo động giả hơn) để đảm bảo an toàn hệ thống.

12 TUNING THRESHOLD

```
[28]: # Lấy probabilities
y_prob_final_xgb = final_pipe_xgb.predict_proba(X_test)[: , 1]

# Tìm threshold tối ưu
precisions, recalls, thresholds = precision_recall_curve(y_test,
    ↪ y_prob_final_xgb)
f1_scores = 2 * (precisions[:-1] * recalls[:-1]) / (precisions[:-1] + recalls[:
    ↪ -1] + 1e-10)

# Threshold với F1 cao nhất
best_idx = np.argmax(f1_scores)
best_threshold_xgb = thresholds[best_idx]
best_f1_xgb = f1_scores[best_idx]

# Threshold với recall cao ( $\geq 0.9$ ) và precision tốt nhất
high_recall_indices = np.where(recalls[:-1]  $\geq$  0.9)[0]
if len(high_recall_indices) > 0:
    best_high_recall_idx = high_recall_indices[np.
    ↪ argmax(precisions[high_recall_indices])]
    high_recall_threshold_xgb = thresholds[best_high_recall_idx]
else:
    high_recall_threshold_xgb = best_threshold_xgb

print(f"Best threshold (max F1): {best_threshold_xgb:.4f} (F1 = {best_f1_xgb:.
    ↪ 4f})")
print(f"High-recall threshold (recall  $\geq$  0.9): {high_recall_threshold_xgb:.4f}")

# Dùng threshold tối ưu
y_pred_final_xgb = (y_prob_final_xgb  $\geq$  best_threshold_xgb).astype(int)
```

Best threshold (max F1): 0.8257 (F1 = 0.7255)

High-recall threshold (recall $\geq$ 0.9): 0.1974

13 ĐÁNH GIÁ MÔ HÌNH CUỐI

```
[29]: # Tính metrics
roc_auc_final_xgb = roc_auc_score(y_test, y_prob_final_xgb)
pr_auc_final_xgb = average_precision_score(y_test, y_prob_final_xgb)
recall_final_xgb = recall_score(y_test, y_pred_final_xgb)
precision_final_xgb = precision_score(y_test, y_pred_final_xgb)
f1_final_xgb = f1_score(y_test, y_pred_final_xgb)

print("Classification Report:")
print(classification_report(y_test, y_pred_final_xgb, digits=4))

print(f"\nROC-AUC: {roc_auc_final_xgb:.6f}")
print(f"PR-AUC: {pr_auc_final_xgb:.6f}")
print(f"Recall: {recall_final_xgb:.4f}")
print(f"Precision: {precision_final_xgb:.4f}")
print(f"F1-Score: {f1_final_xgb:.4f}")

print("\nConfusion Matrix:")
cm_final_xgb = confusion_matrix(y_test, y_pred_final_xgb)
print(cm_final_xgb)
```

Classification Report:

	precision	recall	f1-score	support
0	0.9910	0.9963	0.9937	63544
1	0.8198	0.6506	0.7255	1643
accuracy			0.9876	65187
macro avg	0.9054	0.8235	0.8596	65187
weighted avg	0.9867	0.9876	0.9869	65187

ROC-AUC: 0.971583

PR-AUC: 0.763503

Recall: 0.6506

Precision: 0.8198

F1-Score: 0.7255

Confusion Matrix:

```
[[63309  235]
 [ 574 1069]]
```

14 SO SÁNH KẾT QUẢ

```
[30]: # Tính metrics baseline (không resampling, với scale_pos_weight)
roc_auc_baseline_xgb = roc_auc_score(y_test, y_prob_baseline)
pr_auc_baseline_xgb = average_precision_score(y_test, y_prob_baseline)
recall_baseline_xgb = recall_score(y_test, y_pred_baseline)
precision_baseline_xgb = precision_score(y_test, y_pred_baseline)
f1_baseline_xgb = f1_score(y_test, y_pred_baseline)

print("\nBaseline (scale_pos_weight only):")
print(f" ROC-AUC: {roc_auc_baseline_xgb:.6f}")
print(f" PR-AUC: {pr_auc_baseline_xgb:.6f}")
print(f" Recall: {recall_baseline_xgb:.4f}")
print(f" Precision: {precision_baseline_xgb:.4f}")
print(f" F1-Score: {f1_baseline_xgb:.4f}")

print(f"\nFinal Model ({best_strategy_xgb.upper()} + Tuning):")
print(f" ROC-AUC: {roc_auc_final_xgb:.6f}")
print(f" PR-AUC: {pr_auc_final_xgb:.6f}")
print(f" Recall: {recall_final_xgb:.4f}")
print(f" Precision: {precision_final_xgb:.4f}")
print(f" F1-Score: {f1_final_xgb:.4f}")
```

Baseline (scale_pos_weight only):

ROC-AUC: 0.975064
PR-AUC: 0.788306
Recall: 0.8716
Precision: 0.2830
F1-Score: 0.4273

Final Model (SMOTE + Tuning):

ROC-AUC: 0.971583
PR-AUC: 0.763503
Recall: 0.6506
Precision: 0.8198
F1-Score: 0.7255

```
[31]: # So sánh chi tiết
comparison_xgb = pd.DataFrame({
    'Metric': ['ROC-AUC', 'PR-AUC', 'Recall', 'Precision', 'F1-Score'],
    'Baseline (scale_pos_weight)': [
        f'{roc_auc_baseline_xgb:.6f}',
        f'{pr_auc_baseline_xgb:.6f}',
        f'{recall_baseline_xgb:.4f}',
        f'{precision_baseline_xgb:.4f}',
        f'{f1_baseline_xgb:.4f}'
    ]
})
```

```

],
f'With {best_strategy_xgb.upper()} + Tuning': [
    f'{roc_auc_final_xgb:.6f}',
    f'{pr_auc_final_xgb:.6f}',
    f'{recall_final_xgb:.4f}',
    f'{precision_final_xgb:.4f}',
    f'{f1_final_xgb:.4f}'
],
'Improvement': [
    f'{{{roc_auc_final_xgb - roc_auc_baseline_xgb} / roc_auc_baseline_xgb * 100}:.2f}%',
    f'{{{pr_auc_final_xgb - pr_auc_baseline_xgb} / pr_auc_baseline_xgb * 100}:.2f}%',
    f'{{{recall_final_xgb - recall_baseline_xgb} / recall_baseline_xgb * 100}:.2f}%',
    f'{{{precision_final_xgb - precision_baseline_xgb} / precision_baseline_xgb * 100}:.2f}%',
    f'{{{f1_final_xgb - f1_baseline_xgb} / f1_baseline_xgb * 100}:.2f}%'
]
})

print("\nSo sánh chi tiết:")
print(comparison_xgb.to_string(index=False))

```

So sánh chi tiết:

Metric	Baseline (scale_pos_weight)	With SMOTE + Tuning	Improvement
ROC-AUC	0.975064	0.971583	-0.36%
PR-AUC	0.788306	0.763503	-3.15%
Recall	0.8716	0.6506	-25.35%
Precision	0.2830	0.8198	189.67%
F1-Score	0.4273	0.7255	69.79%

15 VẼ BIỂU ĐỒ

```

[32]: fig, axes = plt.subplots(2, 3, figsize=(18, 12))

# 1. ROC Curve
fpr_baseline, tpr_baseline, _ = roc_curve(y_test, y_prob_baseline)
fpr_final, tpr_final, _ = roc_curve(y_test, y_prob_final_xgb)

axes[0, 0].plot(fpr_baseline, tpr_baseline, label=f'Baseline (AUC = {roc_auc_baseline_xgb:.3f})', linewidth=2)
axes[0, 0].plot(fpr_final, tpr_final, label=f'Final (AUC = {roc_auc_final_xgb:.3f})', linewidth=2)
axes[0, 0].plot([0, 1], [0, 1], 'k--', alpha=0.5)

```

```

axes[0, 0].set_xlabel('False Positive Rate', fontsize=12)
axes[0, 0].set_ylabel('True Positive Rate', fontsize=12)
axes[0, 0].set_title('ROC Curves Comparison (XGBoost)', fontsize=14,
    ↳fontweight='bold')
axes[0, 0].legend()
axes[0, 0].grid(True, alpha=0.3)

# 2. Precision-Recall Curve
precision_baseline_curve, recall_baseline_curve, _ =
    ↳precision_recall_curve(y_test, y_prob_baseline)
precision_final_curve, recall_final_curve, _ = precision_recall_curve(y_test,
    ↳y_prob_final_xgb)

axes[0, 1].plot(recall_baseline_curve, precision_baseline_curve,
    label=f'Baseline (AP = {pr_auc_baseline_xgb:.3f})', linewidth=2)
axes[0, 1].plot(recall_final_curve, precision_final_curve,
    label=f'Final (AP = {pr_auc_final_xgb:.3f})', linewidth=2)
axes[0, 1].set_xlabel('Recall', fontsize=12)
axes[0, 1].set_ylabel('Precision', fontsize=12)
axes[0, 1].set_title('Precision-Recall Curves (XGBoost)', fontsize=14,
    ↳fontweight='bold')
axes[0, 1].legend()
axes[0, 1].grid(True, alpha=0.3)

# 3. Confusion Matrix (Final)
sns.heatmap(cm_final_xgb, annot=True, fmt='d', cmap='Blues', ax=axes[0, 2])
axes[0, 2].set_xlabel('Predicted', fontsize=12)
axes[0, 2].set_ylabel('Actual', fontsize=12)
axes[0, 2].set_title('Confusion Matrix (Final XGBoost Model)', fontsize=14,
    ↳fontweight='bold')

# 4. Threshold Analysis
axes[1, 0].plot(thresholds, f1_scores, 'b-', linewidth=2)
axes[1, 0].axvline(x=best_threshold_xgb, color='r', linestyle='--',
    ↳label=f'Best threshold = {best_threshold_xgb:.3f}')
axes[1, 0].set_xlabel('Threshold', fontsize=12)
axes[1, 0].set_ylabel('F1-Score', fontsize=12)
axes[1, 0].set_title('Threshold vs F1-Score (XGBoost)', fontsize=14,
    ↳fontweight='bold')
axes[1, 0].legend()
axes[1, 0].grid(True, alpha=0.3)

# 5. Feature Importance
feature_names_xgb = final_pipe_xgb.named_steps['preprocess'].
    ↳get_feature_names_out()
importances_xgb = final_pipe_xgb.named_steps['clf'].feature_importances_

```

```

feature_importance_df_xgb = pd.DataFrame({
    'feature': feature_names_xgb,
    'importance': importances_xgb
}).sort_values('importance', ascending=False).head(15)


axes[1, 1].barh(range(len(feature_importance_df_xgb)),
    ↪feature_importance_df_xgb['importance'])
axes[1, 1].set_yticks(range(len(feature_importance_df_xgb)))
axes[1, 1].set_yticklabels(feature_importance_df_xgb['feature'])
axes[1, 1].set_xlabel('Importance', fontsize=12)
axes[1, 1].set_title('Top 15 Feature Importances (XGBoost)', fontsize=14,
    ↪fontweight='bold')
axes[1, 1].grid(True, alpha=0.3, axis='x')

# 6. Comparison Bar Chart
metrics = ['ROC-AUC', 'PR-AUC', 'Recall', 'Precision', 'F1-Score']
baseline_values_xgb = [roc_auc_baseline_xgb, pr_auc_baseline_xgb,
    ↪recall_baseline_xgb, precision_baseline_xgb, f1_baseline_xgb]
final_values_xgb = [roc_auc_final_xgb, pr_auc_final_xgb, recall_final_xgb,
    ↪precision_final_xgb, f1_final_xgb]

x = np.arange(len(metrics))
width = 0.35
axes[1, 2].bar(x - width/2, baseline_values_xgb, width, label='Baseline',
    ↪alpha=0.8)
axes[1, 2].bar(x + width/2, final_values_xgb, width, label='Final', alpha=0.8)
axes[1, 2].set_xlabel('Metrics', fontsize=12)
axes[1, 2].set_ylabel('Score', fontsize=12)
axes[1, 2].set_title('Performance Comparison (XGBoost)', fontsize=14,
    ↪fontweight='bold')
axes[1, 2].set_xticks(x)
axes[1, 2].set_xticklabels(metrics, rotation=45, ha='right')
axes[1, 2].legend()
axes[1, 2].grid(True, alpha=0.3, axis='y')

plt.tight_layout()
plt.savefig('xgboost_evaluation.png', dpi=300, bbox_inches='tight')
plt.show()

```



16 THỬ NGHIỆM VỚI THRESHOLD KHÁC NHAU

```
[33]: thresholds_to_try = [0.1, 0.2, 0.3, 0.4, 0.5, best_threshold_xgb,
    ↪ high_recall_threshold_xgb]
results_thresholds_xgb = []

for thr in thresholds_to_try:
    y_pred_thr = (y_prob_final_xgb >= thr).astype(int)
    recall_thr = recall_score(y_test, y_pred_thr)
    precision_thr = precision_score(y_test, y_pred_thr)
    f1_thr = f1_score(y_test, y_pred_thr)

    results_thresholds_xgb.append({
        'threshold': thr,
        'recall': recall_thr,
        'precision': precision_thr,
        'f1': f1_thr
    })

    print(f"Threshold {thr:.3f}: Recall = {recall_thr:.4f}, Precision =
    ↪ {precision_thr:.4f}, F1 = {f1_thr:.4f}")
```

```
# Chuyển sang DataFrame
threshold_results_df_xgb = pd.DataFrame(results_thresholds_xgb)
print("\nBảng so sánh thresholds:")
print(threshold_results_df_xgb.to_string(index=False))
```

```
Threshold 0.100: Recall = 0.9416, Precision = 0.1353, F1 = 0.2365
Threshold 0.200: Recall = 0.8990, Precision = 0.2047, F1 = 0.3334
Threshold 0.300: Recall = 0.8472, Precision = 0.3027, F1 = 0.4460
Threshold 0.400: Recall = 0.8089, Precision = 0.3857, F1 = 0.5223
Threshold 0.500: Recall = 0.7742, Precision = 0.4961, F1 = 0.6047
Threshold 0.826: Recall = 0.6506, Precision = 0.8198, F1 = 0.7255
Threshold 0.197: Recall = 0.9002, Precision = 0.2024, F1 = 0.3305
```

Bảng so sánh thresholds:

threshold	recall	precision	f1
0.100000	0.941570	0.135263	0.236544
0.200000	0.898965	0.204656	0.333409
0.300000	0.847231	0.302674	0.446011
0.400000	0.808886	0.385665	0.522303
0.500000	0.774194	0.496100	0.604706
0.825725	0.650639	0.819785	0.725484
0.197357	0.900183	0.202381	0.330466

17 XUẤT BÁO CÁO TÓM TẮT

```
[34]: # So sánh với Random Forest (nếu có)
try:
    rf_eval = joblib.load('C:/Users/hi/Downloads/fraud_detection_project/models/
    rf_evaluation_results.pkl')
    rf_f1 = rf_eval['f1']
    comparison_rf_xgb = f"""
SO SÁNH XGBOOST VS RANDOM FOREST:
---
- Random Forest F1-Score: {rf_f1:.4f}
- XGBoost F1-Score: {f1_final_xgb:.4f}
- Chênh lệch: {((f1_final_xgb - rf_f1) / rf_f1 * 100):.2f}%
"""
    print(comparison_rf_xgb)
except:
    print("Không tìm thấy kết quả Random Forest để so sánh")

print(f"""
TÓM TẮT KẾT QUẢ XGBOOST:
---
1. Chiến lược resampling tốt nhất: {best_strategy_xgb.upper()}
2. Số lượng cây (n_estimators): {best_params_xgb.get('n_estimators', 300)}
""")
```

3. Độ sâu tối đa (max_depth): {best_params_xgb.get('max_depth', 6)}
4. Learning rate: {best_params_xgb.get('learning_rate', 0.1):.4f}
5. Threshold tối ưu: {best_threshold_xgb:.4f}
6. Threshold cho recall cao: {high_recall_threshold_xgb:.4f}

KẾT QUẢ ĐÁNH GIÁ:

- ROC-AUC: {roc_auc_final_xgb:.6f} ({((roc_auc_final_xgb -
↪roc_auc_baseline_xgb)/roc_auc_baseline_xgb*100):+.2f}% so với baseline)
- PR-AUC: {pr_auc_final_xgb:.6f} ({((pr_auc_final_xgb - pr_auc_baseline_xgb)/
↪pr_auc_baseline_xgb*100):+.2f}% so với baseline)
- Recall: {recall_final_xgb:.4f} ({((recall_final_xgb - recall_baseline_xgb)/
↪recall_baseline_xgb*100):+.2f}% so với baseline)
- Precision: {precision_final_xgb:.4f} ({((precision_final_xgb -
↪precision_baseline_xgb)/precision_baseline_xgb*100):+.2f}% so với baseline)
- F1-Score: {f1_final_xgb:.4f} ({((f1_final_xgb - f1_baseline_xgb)/
↪f1_baseline_xgb*100):+.2f}% so với baseline)

FEATURES QUAN TRỌNG NHẤT:

```
{feature_importance_df_xgb.head(5).to_string(index=False)}
```

KHUYẾN NGHỊ:

1. Nếu ưu tiên phát hiện nhiều gian lận (recall cao):
 - Dừng threshold: {high_recall_threshold_xgb:.4f}
 - Recall đạt: ~{threshold_results_df_xgb.
↪loc[threshold_results_df_xgb['threshold'] == high_recall_threshold_xgb,
↪'recall'].values[0]:.2%}
 2. Nếu ưu tiên độ chính xác (precision cao):
 - Dừng threshold: {best_threshold_xgb:.4f}
 - Precision đạt: ~{precision_final_xgb:.2%}
 - F1-Score tối ưu: {best_f1_xgb:.4f}
- """

SO SÁNH XGBOOST VS RANDOM FOREST:

- Random Forest F1-Score: 0.6966
- XGBoost F1-Score: 0.7255
- Chênh lệch: 4.15%

TÓM TẮT KẾT QUẢ XGBOOST:

1. Chiến lược resampling tốt nhất: SMOTE

2. Số lượng cây (n_estimators): 329
3. Độ sâu tối đa (max_depth): 7
4. Learning rate: 0.0306
5. Threshold tối ưu: 0.8257
6. Threshold cho recall cao: 0.1974

KẾT QUẢ ĐÁNH GIÁ:

- ROC-AUC: 0.971583 (-0.36% so với baseline)
- PR-AUC: 0.763503 (-3.15% so với baseline)
- Recall: 0.6506 (-25.35% so với baseline)
- Precision: 0.8198 (+189.67% so với baseline)
- F1-Score: 0.7255 (+69.79% so với baseline)

FEATURES QUAN TRỌNG NHẤT:

feature	importance
num__isCustomerDest	0.233105
cat__type_PAYMENT	0.181995
cat__type_TRANSFER	0.178822
num__isMerchant	0.137429
cat__type_CASH_OUT	0.102855

KHUYẾN NGHỊ:

1. Nếu ưu tiên phát hiện nhiều gian lận (recall cao):
 - Dừng threshold: 0.1974
 - Recall đạt: ~90.02%
2. Nếu ưu tiên độ chính xác (precision cao):
 - Dừng threshold: 0.8257
 - Precision đạt: ~81.98%
 - F1-Score tối ưu: 0.7255

18 LƯU MÔ HÌNH VÀ KẾT QUẢ

```
[35]: # Tạo thư mục
os.makedirs('models', exist_ok=True)

# Lưu model
model_path_xgb = 'C:/Users/hi/Downloads/fraud_detection_project/models/
↳xgboost_final.pkl'
joblib.dump(final_pipe_xgb, model_path_xgb)
print(f" Model đã được lưu tại: {model_path_xgb}")
```

```

# Lưu threshold
threshold_data_xgb = {
    'best_threshold': best_threshold_xgb,
    'high_recall_threshold': high_recall_threshold_xgb,
    'best_f1': best_f1_xgb
}
threshold_path_xgb = 'C:/Users/hi/Downloads/fraud_detection_project/models/
↳xgb_best_threshold.pkl'
joblib.dump(threshold_data_xgb, threshold_path_xgb)
print(f" Threshold data đã được lưu tại: {threshold_path_xgb}")

# Lưu feature importance
feature_importance_path_xgb = 'C:/Users/hi/Downloads/fraud_detection_project/
↳models/xgb_feature_importance.csv'
feature_importance_df_xgb.to_csv(feature_importance_path_xgb, index=False)
print(f" Feature importance đã được lưu tại: {feature_importance_path_xgb}")

# Lưu kết quả đánh giá
eval_results_xgb = {
    'roc_auc': roc_auc_final_xgb,
    'pr_auc': pr_auc_final_xgb,
    'recall': recall_final_xgb,
    'precision': precision_final_xgb,
    'f1': f1_final_xgb,
    'confusion_matrix': cm_final_xgb.tolist()
}
eval_path_xgb = 'C:/Users/hi/Downloads/fraud_detection_project/models/
↳xgb_evaluation_results.pkl'
joblib.dump(eval_results_xgb, eval_path_xgb)
print(f" Evaluation results đã được lưu tại: {eval_path_xgb}")

# Lưu best parameters
params_path_xgb = 'C:/Users/hi/Downloads/fraud_detection_project/models/
↳xgb_best_params.pkl'
joblib.dump(best_params_xgb, params_path_xgb)
print(f" Best parameters đã được lưu tại: {params_path_xgb}")

```

Model đã được lưu tại:
C:/Users/hi/Downloads/fraud_detection_project/models/xgboost_final.pkl
Threshold data đã được lưu tại:
C:/Users/hi/Downloads/fraud_detection_project/models/xgb_best_threshold.pkl
Feature importance đã được lưu tại:
C:/Users/hi/Downloads/fraud_detection_project/models/xgb_feature_importance.csv
Evaluation results đã được lưu tại:
C:/Users/hi/Downloads/fraud_detection_project/models/xgb_evaluation_results.pkl
Best parameters đã được lưu tại:
C:/Users/hi/Downloads/fraud_detection_project/models/xgb_best_params.pkl

18.0.1 TỔNG KẾT DỰ ÁN (FINAL VERDICT)

Sau khi thử nghiệm 3 mô hình, chúng ta có bảng so sánh hiệu suất trên tập Test (Hold-out set):

Metric	Logistic Regression	Random Forest	XGBoost (Champion)
Resampling	SMOTE+ENN	SMOTE	SMOTE
ROC-AUC	0.9667	0.9664	0.9716
Precision	0.5707	0.7682	0.8198
Recall	0.5648	0.6372	0.6506
F1-Score	0.5678	0.6966	0.7255

Key Takeaways:

1. **XGBoost là mô hình vượt trội nhất:** F1-Score đạt **72.55%**, cao hơn Random Forest ~4% và Logistic Regression ~16%.
2. **Threshold tối ưu:** Nên sử dụng threshold **0.8257** để đạt hiệu suất cao nhất. Nếu cần an toàn tuyệt đối, dùng threshold **0.1974** (Recall > 90%).
3. **Đặc trưng hành vi:** Gian lận tập trung vào **TRANSFER/CASH_OUT** và có xu hướng theo khung giờ (hour).

Khuyến nghị: Triển khai **XGBoost** cho hệ thống Production.